

a glance at reinforcement learning

adam green

email - adam.green@adagefficiency.com slides - notes - practical work -



agenda

about me

goals for this course

where to go next



expectations & conditionals

variance & bias

bootstrapping

function approximation

iid



learning rate

batch size

scaling / preprocessing

batch norm



applications (maybe should go before context)

context



value functions & DQN (TODO expand)



practical - just go to the readme -> better readme



dqn extensions

pg

alpha go

practical stuff

state of the art

